off	h → shared regression head
simple_concat	[h, c] → Linear (464 → 400) → shared head
concat	[h, c] → Linear → ReLU → LayerNorm → shared head
film	$h' = (1 + g(c)) \odot h + b(c)$ → shared head

h: 400-d fingerprint; c: 64-d condition code. FiLM projection weights and biases start at zero.